

# VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



**LAB REPORT**  
**on**

## **Machine Learning (23CS6PCMAL)**

*Submitted by*

**NITHYA LAKSHMI V (1BM22CS186)**

*in partial fulfillment for the award of the degree of*

**BACHELOR OF ENGINEERING**  
*in*  
**COMPUTER SCIENCE AND ENGINEERING**



**B.M.S. COLLEGE OF ENGINEERING**

(Autonomous Institution under VTU)

**BENGALURU-560019**

**Sep-2024 to Jan-2025**

**B.M.S. College of Engineering,**  
**Bull Temple Road, Bangalore 560019**  
(Affiliated To Visvesvaraya Technological University, Belgaum)  
**Department of Computer Science and Engineering**



**CERTIFICATE**

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **NITHYA LAKSHMI V (1BM22CS186)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Ms. Saritha A N Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
--	--

## Index

Sl. No.	Date	Experiment Title	Page No.
1	3-3-2025	Write a python program to import and export data using Pandas library functions	1 - 3
2	10/03/25	Demonstrate various data pre-processing techniques for a given dataset	4 - 8
3	10-3-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	9 - 15
4	17-3-2025	Build Logistic Regression Model for a given dataset	16 - 18
5	24-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	19 - 21
6	7-4-2025	Build KNN Classification model for a given dataset.	22 - 25
7	21-4-2025	Build Support vector machine model for a given dataset	26 - 30
8	5-5-2025	Implement Random forest ensemble method on a given dataset.	31 - 33
9	5-5-2025	Implement Boosting ensemble method on a given dataset.	34 - 39
10	05/05/25	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	40 - 45
11	05/05/25	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	46 - 50

Github Link:

<https://github.com/Nithya1909/ML-LAB>

### **Program 1**

Write a python program to import and export data using Pandas library functions

Code:

```
import pandas as pd

# Method-1: Initializing values directly into DataFrame

data_method1 = {'USN': ['1JS17CS001', '1JS17CS002', '1JS17CS003',
                        '1JS17CS004', '1JS17CS005'],
                'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],
                'Marks': [90, 85, 92, 78, 88]}

df_method1 = pd.DataFrame(data_method1)

print("Method-1:")

print(df_method1)

print("-" * 20)

# Method-2: Importing datasets from sklearn.datasets

from sklearn.datasets import load_diabetes

diabetes_data = load_diabetes()

df_method2 = pd.DataFrame(data=diabetes_data.data,
                          columns=diabetes_data.feature_names)

df_method2['target'] = diabetes_data.target

print("Method-2:")

print(df_method2.head())

print("-" * 20)

# Method-3: Importing datasets from a specific .csv file

try:
```

```

df_method3 = pd.read_csv('sample_sales_data.csv')

print("Method-3:")

print(df_method3.head())

print("-" * 20)

except FileNotFoundError:

print("sample_sales_data.csv not found. Please upload the file.")

print("-" * 20)

import yfinance as yf

import matplotlib.pyplot as plt

tickers = ["HDFCBANK.NS", "ICICIBANK.NS", "KOTAKBANK.NS"]

start_date = "2024-01-01"

end_date = "2024-12-30"

data = yf.download(tickers, start=start_date, end=end_date)

closing_prices = data['Close']

daily_returns = closing_prices.pct_change().dropna()

plt.figure(figsize=(12, 6))

closing_prices.plot()

plt.title('Closing Prices (2024)')

plt.xlabel('Date')

plt.ylabel('Price (INR)')

plt.grid(True)

plt.show()

plt.figure(figsize=(12, 6))

daily_returns.plot()

plt.title('Daily Returns (2024)')

plt.xlabel('Date')

```

```
plt.ylabel('Daily Return')
```

```
plt.grid(True)
```

```
plt.show()
```

## **Program 2**

Demonstrate various data pre-processing techniques for a given dataset.

Code:

```
from google.colab import files

uploaded = files.upload()

uploaded = files.upload()

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.impute import SimpleImputer

from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder

from sklearn.preprocessing import StandardScaler, MinMaxScaler

from scipy import stats

import pandas as pd
df1=pd.read_csv("adult.csv")

print(df1.head())

df2=pd.read_csv("Dataset of Diabetes .csv")

print(df2.head())

df1["education"].value_counts()

ordinal_encoder = OrdinalEncoder(categories=[["HS-grad",
"Some-college", "Bachelors", "Masters", "Assoc-voc", "11th", "Assoc-acdm", "10th", "7t
h-8th", "Prof-school", "9th", "12th", "Doctorate", "5th-6th", "1st-4th", "Preschool"]])

df1["Education_Encoded"] = ordinal_encoder.fit_transform(df1[["education"]])

onehot_encoder = OneHotEncoder()

encoded_data =

onehot_encoder.fit_transform(df1[["gender", "relationship", "workclass", "occupati
on", "race", "native-country", "income", "marital-status"]])
```

```

encoded_array = encoded_data.toarray()

encoded_df = pd.DataFrame(encoded_array,
columns=onehot_encoder.get_feature_names_out(["gender","relationship","workclas
s","occupation","race","native-country","income","marital-status"]))

df_encoded = pd.concat([df1, encoded_df], axis=1)

df_encoded.drop(["education","gender","workclass","relationship","occupation","
race","native-country","income","marital-status"], axis=1, inplace=True)

print(df_encoded.head())

normalizer = MinMaxScaler()

df_encoded[["fnlwgt","educational-num","capital-gain","capital-loss","hours-per
-week"]] =
normalizer.fit_transform(df_encoded[["fnlwgt","educational-num","capital-gain",
"capital-loss","hours-per-week"]])

df_encoded.head()

df2.isnull().sum()

df2['Gender'] = df2['Gender'].replace('f', 'F')

ordinal_encoder = OrdinalEncoder(categories=[["F", "M"]])

df2["Gender_Encoded"] = ordinal_encoder.fit_transform(df2[["Gender"]])

onehot_encoder = OneHotEncoder()

encoded_data = onehot_encoder.fit_transform(df2[["CLASS"]])

encoded_array = encoded_data.toarray()

encoded_df = pd.DataFrame(encoded_array,
columns=onehot_encoder.get_feature_names_out(["CLASS"]))

df_encoded = pd.concat([df2, encoded_df], axis=1)

```



```

df2 = pd.concat([df2, encoded_df], axis=1)

df2.drop("CLASS", axis=1, inplace=True)

df2.drop("Gender", axis=1, inplace=True)

print(df2.head())

normalizer = MinMaxScaler()

df_encoded[["No_Pation", "AGE", "Urea", "Cr", "HbA1c",
            "Chol", "TG", "HDL", "LDL", "VLDL", "BMI"]] =
normalizer.fit_transform(df_encoded[["No_Pation", "AGE", "Urea", "Cr", "HbA1c",
            "Chol", "TG", "HDL", "LDL", "VLDL", "BMI"]])

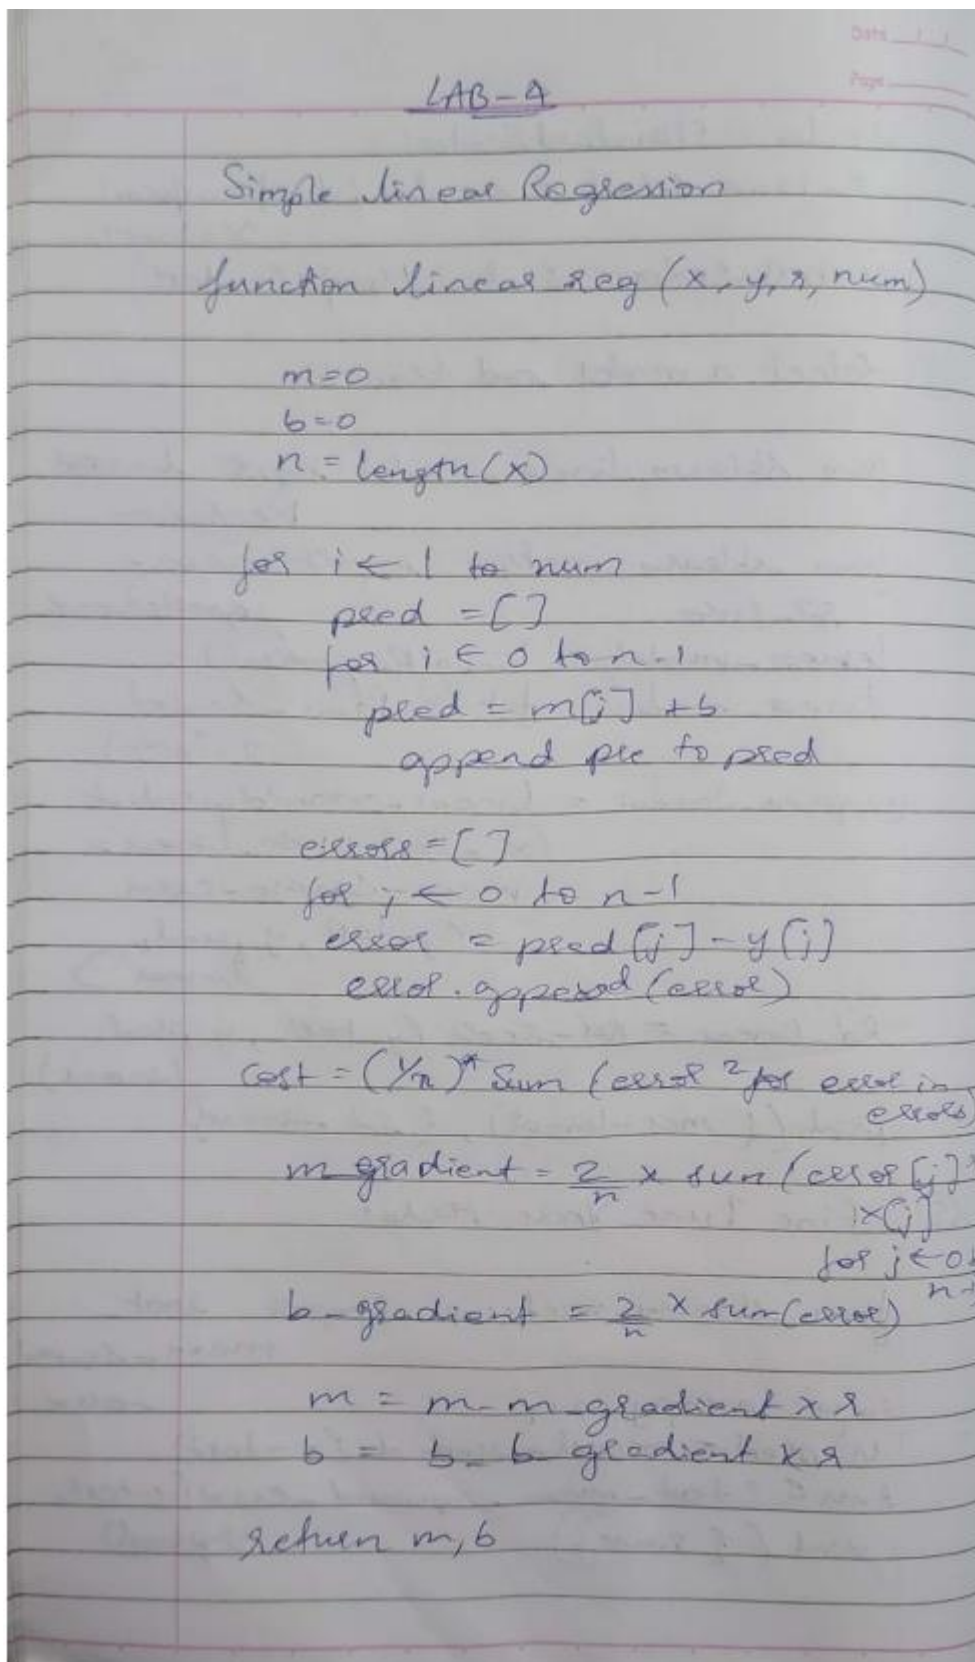
df_encoded.head()

```

### Program 3

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot:



The image shows a handwritten MATLAB script for Simple Linear Regression on lined paper. The title 'LAB-4' is written at the top. The code defines a function 'linear reg' that takes inputs x, y, x, and num. It initializes m=0, b=0, and n=length(x). It then iterates through the data to calculate predicted values and errors. Finally, it calculates the cost function and the gradients for m and b, and returns the updated values.

```
LAB-4

Simple Linear Regression

function linear reg (x, y, x, num)

    m=0
    b=0
    n = length(x)

    for i = 1 to num
        pred = []
        for i = 0 to n-1
            pred = m[i] * x + b
            append pre to pred
        endfor
    endfor

    errors = []
    for j = 0 to n-1
        error = pred[j] - y(j)
        error.append(error)
    endfor

    cost = (1/n) * sum (error^2 for error in errors)

    m_gradient = 2/n * sum (error[j] * x[j])
    b_gradient = 2/n * sum (error)

    m = m - m_gradient * x
    b = b - b_gradient * x

    return m, b
```

### Algo Linear Regression

1. Initialise  $m$  &  $b = 0$
2. make prediction  $mX + b$
3. Find error by  $y - \text{pred}$
4. update  $m$  &  $b$  using gradient descent
5. Repeat 4 steps for num-iterations

### Algo Logistic Regression

1.  $m$  be the length of features initialize to 0
2.  $n$  = no of data points
3. Initialize weight & bias to 0
4. Make predictions weight  $X + \text{bias}$
5. Use Sigmoid function for pred.
6. Find error using  $\text{pred} - y_i$
7. Update weight & bias using gradient descent
8. Repeat steps 4 return weight & bias

### Algo multivariate Regression

Same as linear but with multiple features.

Code:

```
import pandas as pd

import numpy as np

from sklearn import linear_model

import matplotlib.pyplot as plt

df = pd.read_csv('/content/housing_area_price.csv')

df

plt.xlabel('area')

plt.ylabel('price')

plt.scatter(df.area,df.price,color='red',marker='+')

new_df = df.drop('price',axis='columns')

new_df

price = df.price

price

reg = linear_model.LinearRegression()

reg.fit(new_df,price)

print(reg.coef_)

print(reg.intercept_)

reg.predict([[5000]])

df_mlr=pd.read_csv('/content/homeprices_Multiple_LR.csv')

df_mlr

df_mlr.bedrooms.median()

df_mlr.bedrooms = df_mlr.bedrooms.fillna(df_mlr.bedrooms.median())

print(df_mlr)

reg = linear_model.LinearRegression()

reg.fit(df_mlr.drop('price',axis='columns'),df_mlr.price)
```

```

reg.coef_

reg.intercept_

#Find price of home with 3000 sqr ft area, 3 bedrooms, 40 year old

reg.predict([[3000, 3, 40]])

112.06244194*3000 + 23388.88007794*3 + -3231.71790863*40 + 221323.00186540384


canada=pd.read_csv('/content/canada_per_capita_income.csv')

canada

plt.xlabel('year')

plt.ylabel('per capita income (US$)')

canada.rename(columns={'per capita income (US$)':'income'},inplace=True)

plt.scatter(canada.year,canada.income ,color='red',marker='+')

new_df_canada = canada.drop('income',axis='columns')

new_df_canada.sample(5)

income = canada.income

income.sample(5)

reg = linear_model.LinearRegression()

reg.fit(new_df_canada,income)

print(reg.coef_)

print(reg.intercept_)

reg.predict([[2020]])

hiring=pd.read_csv('/content/hiring.csv')

hiring

hiring['experience'].fillna(0, inplace=True)

hiring['test_score(out of 10)'].fillna(hiring['test_score(out of 10)'].mean(),

inplace=True)

```

```

def convert_to_int(word):

word_dict = {'one':1, 'two':2, 'three':3, 'four':4, 'five':5, 'six':6,
'seven':7, 'eight':8,
'nine':9, 'ten':10, 'eleven':11, 'twelve':12, 'zero':0, 0: 0}

return word_dict[word]

hiring['experience'] = hiring['experience'].apply(lambda x : convert_to_int(x))
hiring

reg = linear_model.LinearRegression()

hiring.rename(columns={'salary($)': 'salary'}, inplace=True)

reg.fit(hiring.drop('salary', axis='columns'), hiring.salary)

print(reg.coef_)

print(reg.intercept_)

#What is the predicted salary for a candidate with 12 years of experience, 10
test score, and 10 interview score?

reg.predict([[12, 10, 10]])


comp=pd.read_csv('/content/1000_Companies.csv')

comp

comp.isnull().sum()

from sklearn.preprocessing import OneHotEncoder

ohe = OneHotEncoder(sparse_output=False, handle_unknown='ignore')

state_encoded = ohe.fit_transform(comp[['State']])

state_encoded_df = pd.DataFrame(state_encoded,

columns=ohe.get_feature_names_out(['State']))

comp = pd.concat([comp, state_encoded_df], axis=1).drop(columns=['State'])

print(comp)

reg = linear_model.LinearRegression()

```

```
reg.fit(comp.drop('Profit',axis='columns'),comp.Profit)

print(reg.coef_)

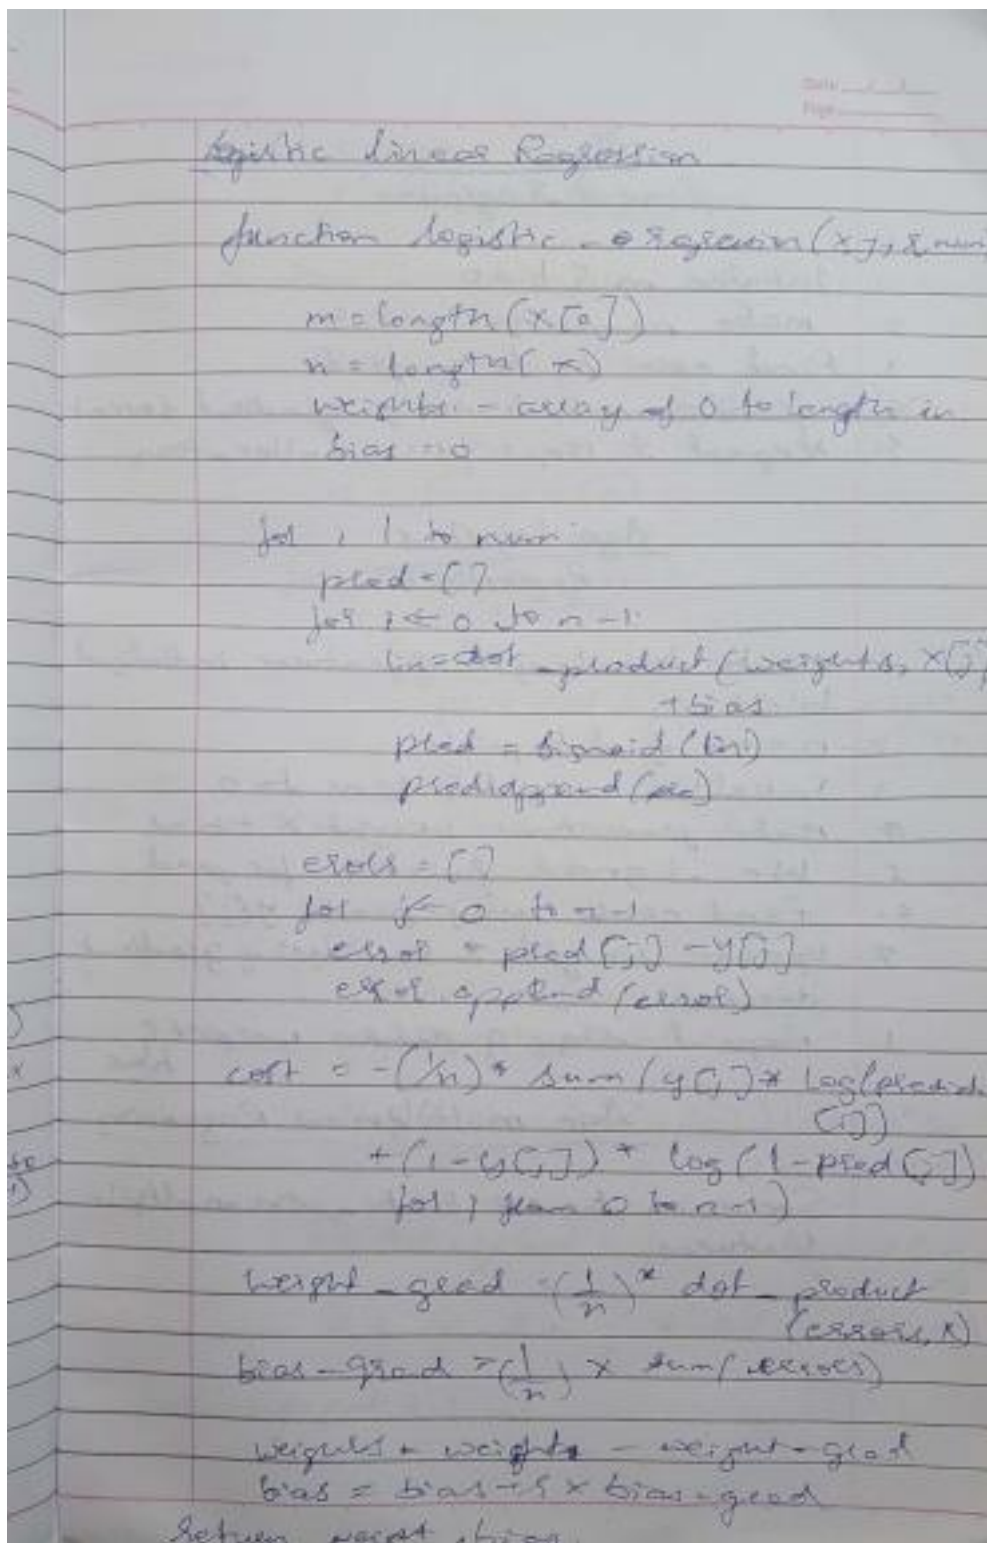
print(reg.intercept_)


reg.predict([[91694.48, 515841.3, 11931.24,0,1,0]])
```

## Program 4

Build Logistic Regression Model for a given dataset.

Screenshot:



```
Logistic Linear Regression

function logistic = @logistic(x, y, num)

m = length(x[0])
n = length(y)
weights = zeros(1, length(x[0]))
bias = 0

for i = 1 to num
    pred = []
    for j = 0 to n-1
        lin = dot_product(weights, x[j]) + bias
        pred = sigmoid(lin)
        pred.append(pred)
    end

    errors = []
    for k = 0 to n-1
        error = pred[k] - y[k]
        error.append(error)
    end

    cost = -(1/n) * sum(y[k] * log(pred[k])
        + (1 - y[k]) * log(1 - pred[k])
        for j from 0 to n-1)

    weight_grad = (1/n) * dot_product(errors, x)
    bias_grad = (1/n) * sum(errors)

    weights = weights - weight_grad
    bias = bias - 1 * bias_grad
end

return weights, bias
```



Code:

```
import pandas as pd

from matplotlib import pyplot as plt

from seaborn import regplot

import seaborn as sns

df1=pd.read_csv('HR_comma_sep.csv')

df1.sample(10)

plt.figure(figsize=(8, 6))

sns.barplot(x='salary_encoded', y='left', data=df1)

plt.title('Impact of Employee Salary on Retention')

plt.xlabel('Salary Level (Encoded)')

plt.ylabel('Proportion of Employees Left')

plt.show()

plt.figure(figsize=(12, 6))

sns.barplot(x='Department', y='left', data=df1)

plt.title('Employee Retention Rate by Department')

plt.xlabel('Department')

plt.ylabel('Proportion of Employees Left')

plt.xticks(rotation=45, ha='right')

plt.show()

from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder

import numpy as np

import seaborn as sns

ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False)

department_encoded = ohe.fit_transform(df1[['Department']])

department_encoded_df = pd.DataFrame(department_encoded,
```

```
columns=ohe.get_feature_names_out(['Department']))

df1 = pd.concat([df1, department_encoded_df], axis=1)

df1 = df1.drop('Department', axis=1)

ordinal_encoder = OrdinalEncoder(categories=[['low', 'medium', 'high']],
dtype=np.int64)

salary_encoded = ordinal_encoder.fit_transform(df1[['salary']])

df1['salary_encoded'] = salary_encoded

df1 = df1.drop('salary', axis=1)

df1.head()

correlation_matrix = df1.corr()

plt.figure(figsize=(12, 10))

sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
```

```

plt.title('Correlation Matrix of Features')

plt.show()

correlation_threshold = 0.1

correlated_features = correlation_matrix['left'].abs() > correlation_threshold

highly_correlated_features =
correlated_features[correlated_features].index.tolist()

new_df = df1[highly_correlated_features]

print(new_df.head())

```

```

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score

```

```

X = new_df.drop('left', axis=1)

y = new_df['left']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

model = LogisticRegression()

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy}")

example_data = X_test.iloc[0].values.reshape(1, -1)

prediction = model.predict(example_data)

```

```

import pandas as pd

```

```

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score

zoo_data = pd.read_csv('zoo-data.csv')

zoo_class = pd.read_csv('zoo-class-type.csv')


merged_data = pd.merge(zoo_data, zoo_class, left_on='class_type',
right_on='Class_Number')

merged_data = merged_data.drop(['Animal_Names',
'Number_Of_Animal_Species_In_Class',
'Class_Number','class_type','animal_name'], axis=1)


X = merged_data.drop('Class_Type', axis=1)
y = merged_data['Class_Type']

print(merged_data.head())


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)


model = LogisticRegression(max_iter=1000)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy}")


from sklearn.metrics import confusion_matrix

import seaborn as sns

```

```
import matplotlib.pyplot as plt

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=np.unique(y_test), yticklabels=np.unique(y_test))

plt.xlabel("Predicted")

plt.ylabel("Actual")

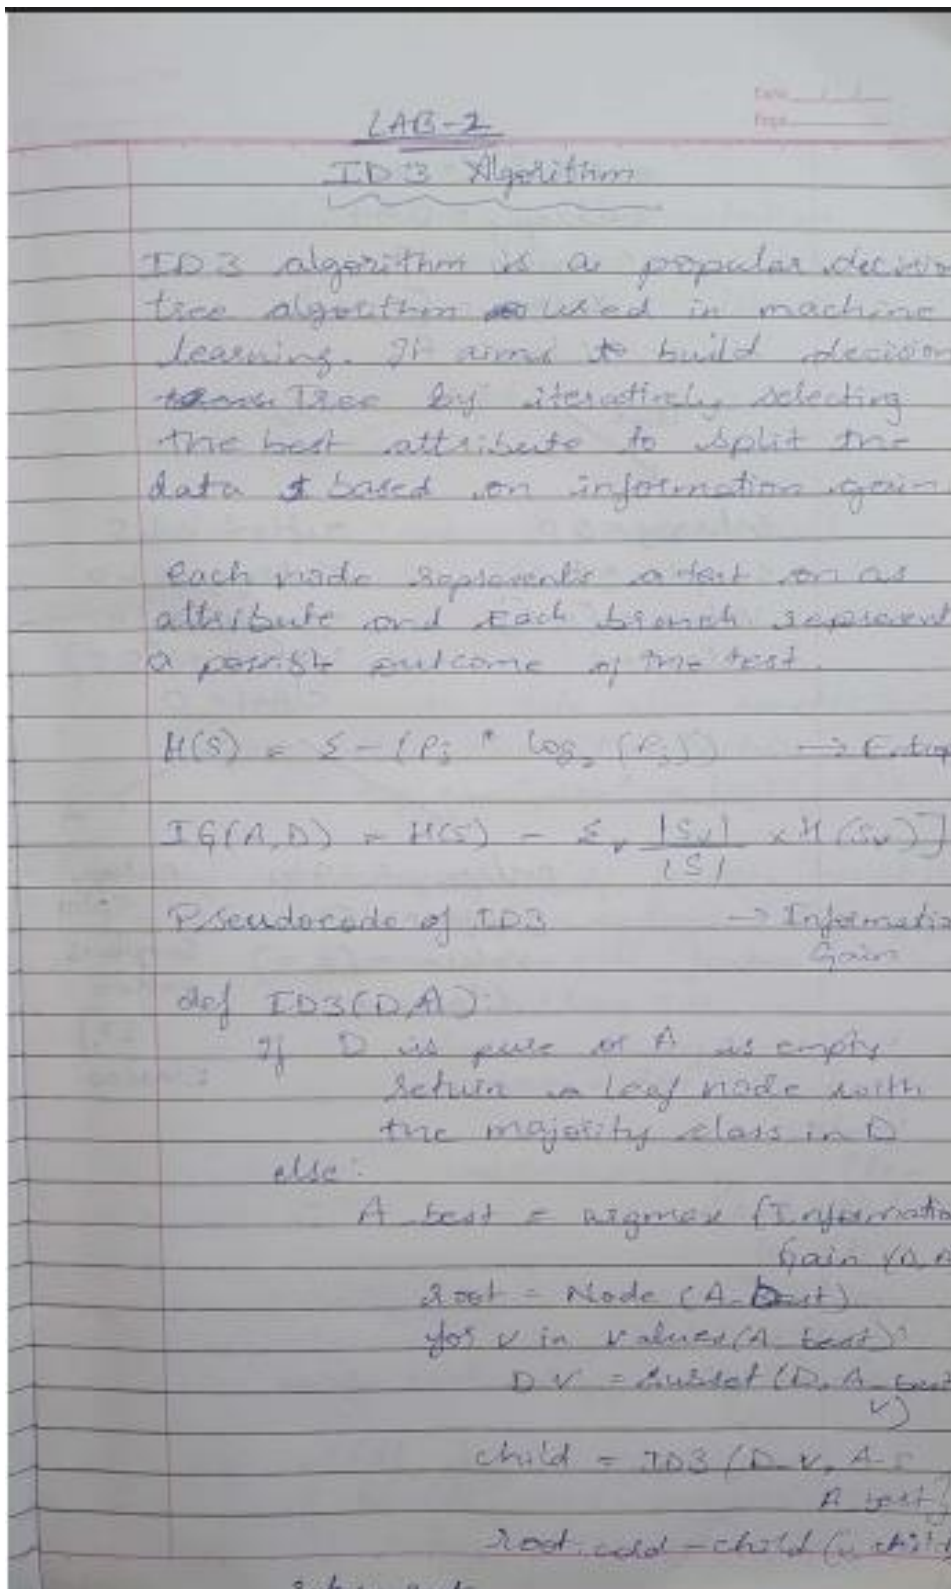
plt.title("Confusion Matrix")

plt.show()
```


## Program 5

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot:




out look  $\leq 10.5$   
 Entropy = 0.94  
 Samples = 14  
 Value = [5, 9]  
 class = 1



Entropy = 0.0  
 Samples = 4  
 Value = [9, 9]  
 class = 1

out look  $\leq 10.5$   
 Entropy = 1.0  
 Samples = 10  
 Value = [5, 9]  
 class = 0



Entropy = 0.971  
 Samples = 5  
 Value = [2, 3]  
 Class = 1

Entropy = 0.917  
 Samples = 5  
 Value = [3, 2]  
 Class = 0

11/03/25

LAB-3

Date: / /  
Page: /

### ID3 - Code Implementation without using sklearn library

```
import pandas as pd
import numpy as np

data = pd.read_csv("weather.csv")
df = pd.DataFrame(data)

def entropy(data):
    class_counts = data.value_counts()
    prob = class_counts / len(data)
    return -np.sum(prob * np.log2(prob))

def information_gain(df, feature, target):
    total_entropy = entropy(df[target])
    feature_value = df[feature].unique()
    weighted_entropy = 0
    for value in feature_value:
        subset = df[df[feature] == value]
        weighted_entropy += (len(subset) / len(df)) * entropy(subset[target])
    return total_entropy - weighted_entropy

def id3(df, features, target):
    if len(df[target].unique()) == 1:
        return df[target].to_numpy()
```



```

if not features:
    return df[target].mode()[0]
gains = [feature_information_gain
          (df, feature, target) for
          feature in features]
best_feature = max(gains, key = gain
                    get)
tree = {best_feature: 1}

for value in df[best_feature].unique():
    subset = df[df[best_feature] == value]
    tree[best_feature][value] = id3(subset,
                                     [f for f in features if
                                      f != best_feature],
                                     target)

return valueficc

```

```

features = ['outlook', 'temperature',
            'humidity', 'wind']
target = 'PlayTennis'
desc_tree = id3(df, features, target)
print(desc_tree)

```

### OUTPUT

```

{ 'Outlook': { 'Sunny':
  { 'Humidity': { 'high': 'No', 'low': 'yes' },
    'outlook': 'Yes', 'Rain': { 'wind':
      { 'weak': 'Yes',
        'Strong': 'No' }}} }

```

Code:

```
import pandas as pd

from sklearn.preprocessing import LabelEncoder

from sklearn.tree import DecisionTreeClassifier

from sklearn.model_selection import train_test_split

from sklearn.metrics import accuracy_score, classification_report

data = {

'a1': [True, True, False, False, False, True, True, True, False, False],

'a2': ['Hot', 'Hot', 'Hot', 'Cool', 'Cool', 'Cool', 'Hot', 'Hot', 'Cool',

'Cool'],

'a3': ['High', 'High', 'High', 'Normal', 'Normal', 'High', 'High',

'Normal', 'Normal', 'High'],

'Classification': ['No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'No', 'Yes',

'Yes', 'Yes']

}

data

df = pd.DataFrame(data)

label_encoders = { }

for column in df.columns:

    le = LabelEncoder()

    df[column] = le.fit_transform(df[column])

    label_encoders[column] = le

df

X = df.drop('Classification', axis=1)

y = df['Classification']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
```

```

random_state=42)

clf = DecisionTreeClassifier(criterion='entropy')

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')

print(classification_report(y_test, y_pred, target_names=['No', 'Yes']))

from sklearn.tree import plot_tree

import matplotlib.pyplot as plt
plt.figure(figsize=(12,8))

plot_tree(clf, filled=True, feature_names=X.columns, class_names=['No', 'Yes'])

plt.show()


iris=pd.read_csv("/content/iris - Copy.csv")

iris

le = LabelEncoder()

iris["species"] = le.fit_transform(iris["species"])

label_encoders[column] = le

iris

X = iris.drop('species', axis=1)

y = iris['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

clf = DecisionTreeClassifier(criterion='entropy')

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

```

```

accuracy = accuracy_score(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')

print(classification_report(y_test, y_pred, target_names=['Iris-setosa',
'Iris-versicolor', 'Iris-virginica',]))

from sklearn.metrics import confusion_matrix

import seaborn as sns

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'],
yticklabels=['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'])

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

plt.show()

plt.figure(figsize=(12,8))

plot_tree(clf, filled=True, feature_names=X.columns,
class_names=['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'])

plt.show()


drug=pd.read_csv("/content/drug - Copy.csv")

print(drug)

drug["Drug"].unique()

le = LabelEncoder()

drug["Sex"] = le.fit_transform(drug["Sex"])

drug["BP"] = le.fit_transform(drug["BP"])

drug["Cholesterol"] = le.fit_transform(drug["Cholesterol"])

```

```

#drug["Drug"] = le.fit_transform(drug["Drug"])

label_encoders[column] = le

drug

X = drug.drop('Drug', axis=1)

y = drug['Drug']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

clf = DecisionTreeClassifier(criterion='entropy')

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')

print(classification_report(y_test, y_pred, target_names=['drugY', 'drugC',
'drugX', 'drugA', 'drugB']))

from sklearn.metrics import confusion_matrix

import seaborn as sns

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=['drugY', 'drugC', 'drugX', 'drugA', 'drugB'],
yticklabels=['drugY', 'drugC', 'drugX', 'drugA', 'drugB'])

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

plt.show()

petrol=pd.read_csv("/content/petrol_consumption - Copy.csv")

```

```

petrol

X = petrol.drop('Petrol_Consumption', axis=1)

y = petrol['Petrol_Consumption']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

clf = DecisionTreeRegressor()

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

from sklearn.tree import DecisionTreeRegressor

from sklearn.model_selection import train_test_split

from sklearn.metrics import mean_absolute_error, mean_squared_error


mae = mean_absolute_error(y_test, y_pred)

mse = mean_squared_error(y_test, y_pred)

rmse = mean_squared_error(y_test, y_pred)**0.5

print(f'Mean Absolute Error: {mae}')

print(f'Mean Squared Error: {mse}')

print(f'Root Mean Squared Error: {rmse}')

plt.figure(figsize=(30, 30))

plot_tree(clf, filled=True, feature_names=X.columns, fontsize=10)

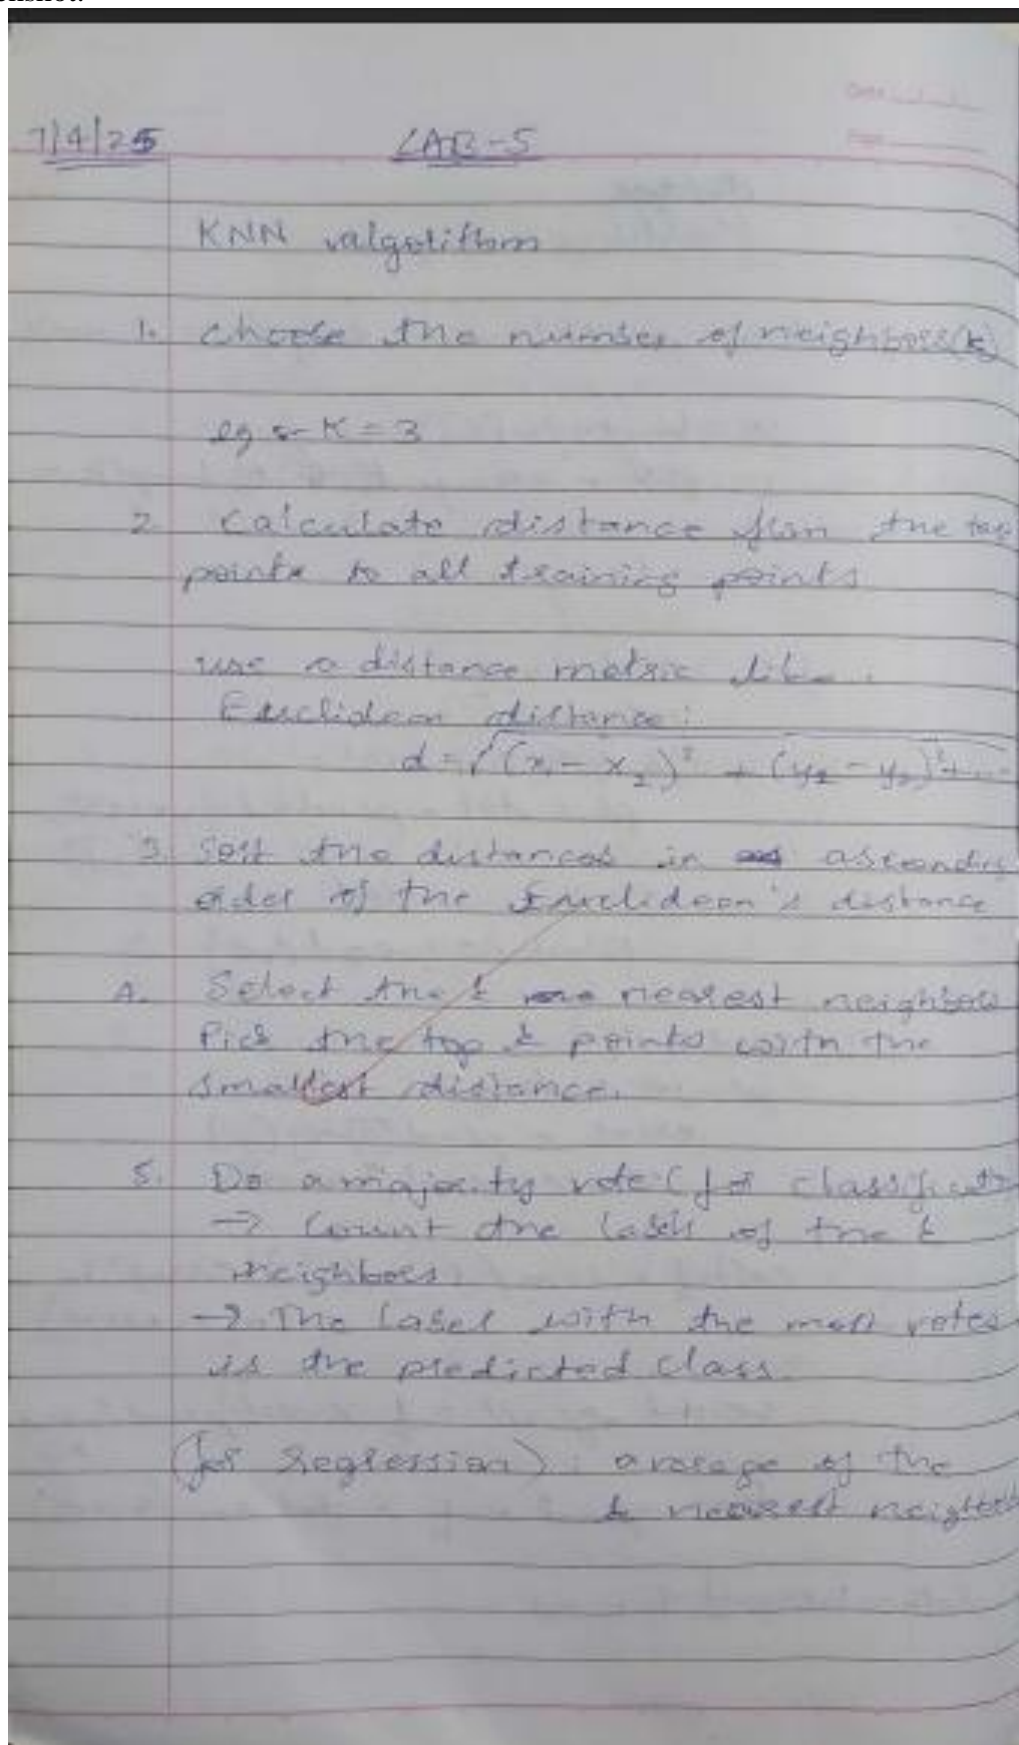
plt.show()

```

## Program 6

Build KNN Classification model for a given dataset.

Screenshot:



Code:

```
import pandas as pd

iris=pd.read_csv('iris.csv')

iris.head()

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, confusion_matrix,

accuracy_score

import matplotlib.pyplot as plt

import seaborn as sns


X = iris.drop('species', axis=1)

y = iris['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,

random_state=42)

best_k = 1

best_accuracy = 0

for k in range(1, 11):

    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X_train, y_train)

    y_pred = knn.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)

    print(f"Accuracy for k={k}: {accuracy}, Error Rate for k={k}:

    { 1-accuracy}")

    if accuracy > best_accuracy:

        best_accuracy = accuracy
```



```

        best_k = k

print(f"Best k value: {best_k}")

knn = KNeighborsClassifier(n_neighbors=best_k)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

print("Accuracy:", accuracy)

cm = confusion_matrix(y_test, y_pred)

print("Confusion Matrix:")

print(cm)

plt.figure(figsize=(8,6))

sns.heatmap(cm, annot=True, fmt='d',
cmap='Blues',xticklabels=iris['species'].unique(),
yticklabels=iris['species'].unique())

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.title("Confusion Matrix")

plt.show()


print("\nClassification Report:")

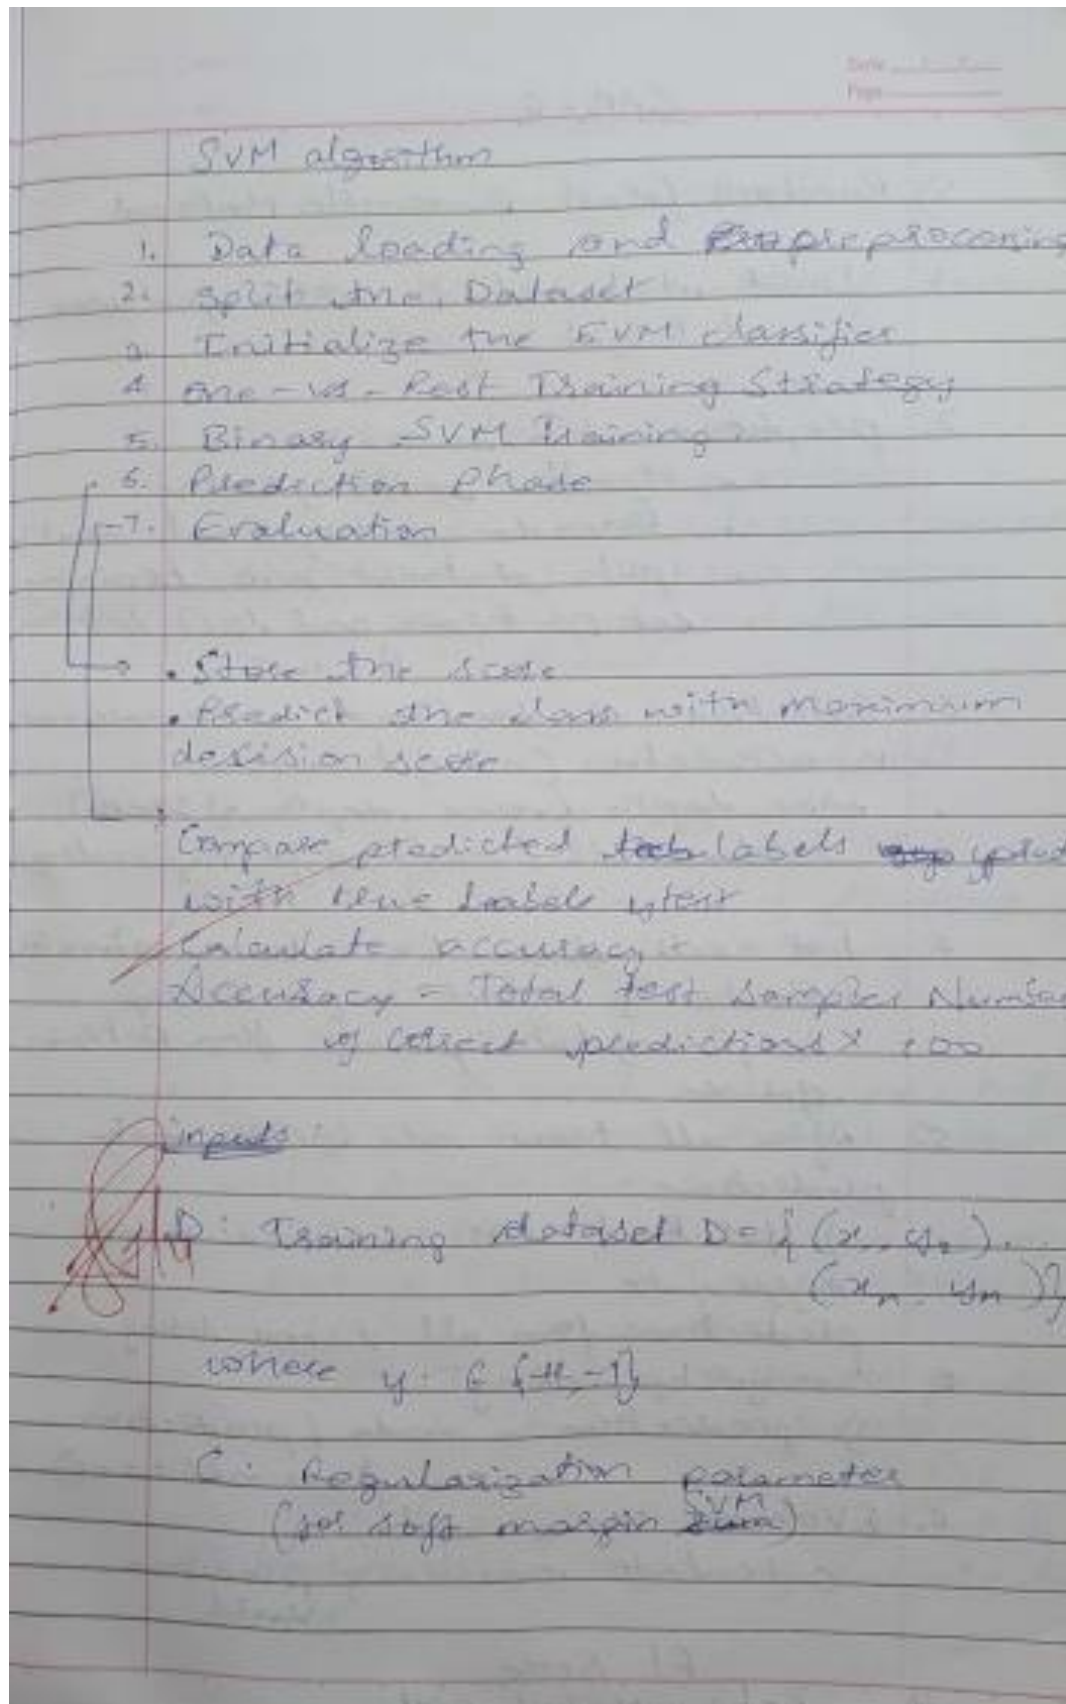
print(classification_report(y_test, y_pred))

```

## Program 7

Build Support vector machine model for a given dataset

Screenshot:



Code:

```
import pandas as pd

iris=pd.read_csv('iris.csv')

iris.head()

iris.isnull().sum()
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, confusion_matrix

import seaborn as sns

import matplotlib.pyplot as plt

iris = pd.read_csv('iris.csv')

X = iris.drop('species', axis=1)

y = iris['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

svm_rbf = SVC(kernel='rbf', gamma='scale')

svm_rbf.fit(X_train, y_train)

y_pred_rbf = svm_rbf.predict(X_test)

accuracy_rbf = accuracy_score(y_test, y_pred_rbf)

print(f"Accuracy (RBF Kernel): { accuracy_rbf}")

cm_rbf = confusion_matrix(y_test, y_pred_rbf)

plt.figure(figsize=(8, 6))

sns.heatmap(cm_rbf, annot=True, fmt="d", cmap="Blues",

xticklabels=iris['species'].unique(),

yticklabels=iris['species'].unique())

plt.title("Confusion Matrix (RBF Kernel)")
```

```

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.show()

svm_linear = SVC(kernel='linear')
svm_linear.fit(X_train, y_train)

y_pred_linear = svm_linear.predict(X_test)

accuracy_linear = accuracy_score(y_test, y_pred_linear)

print(f"Accuracy (Linear Kernel): {accuracy_linear}")

cm_linear = confusion_matrix(y_test, y_pred_linear)

plt.figure(figsize=(8, 6))

sns.heatmap(cm_linear, annot=True, fmt="d", cmap="Blues",
            xticklabels=iris['species'].unique(),
            yticklabels=iris['species'].unique())

plt.title("Confusion Matrix (Linear Kernel)")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.show()

df2=pd.read_csv('letter-recognition.csv')

df2.head()

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

import seaborn as sns

import matplotlib.pyplot as plt

```

```

from sklearn.preprocessing import LabelBinarizer

X = df2.drop('letter', axis=1)

y = df2['letter']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

svm_classifier = SVC(kernel='linear', probability=True)

svm_classifier.fit(X_train, y_train)

y_pred = svm_classifier.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))

print(classification_report(y_test, y_pred))

cm = confusion_matrix(y_test, y_pred)

sns.heatmap(cm, annot=True, fmt='d', cmap="Blues")

plt.title('Confusion Matrix for SVM')

plt.xlabel('Predicted')

plt.ylabel('True')

plt.show()

lb = LabelBinarizer()

lb.fit(y_test)

y_test_lb = lb.transform(y_test)

y_pred_prob = svm_classifier.predict_proba(X_test)

fpr = {}

tpr = {}

thresh = {}

roc_auc = dict()

n_class = y_test_lb.shape[1]

for i in range(n_class):

```

```

fpr[i], tpr[i], thresh[i] = roc_curve(y_test_lb[:,i], y_pred_prob[:,i])

roc_auc[i] = auc(fpr[i], tpr[i])

plt.plot(fpr[0], tpr[0], linestyle='--',color='orange', label='SVM (AUC =
%0.2f)' % roc_auc[0])

plt.title('ROC Curve for Class 0')

plt.xlabel('False Positive Rate')

plt.ylabel('True Positive rate')

plt.legend(loc='best')

plt.show()

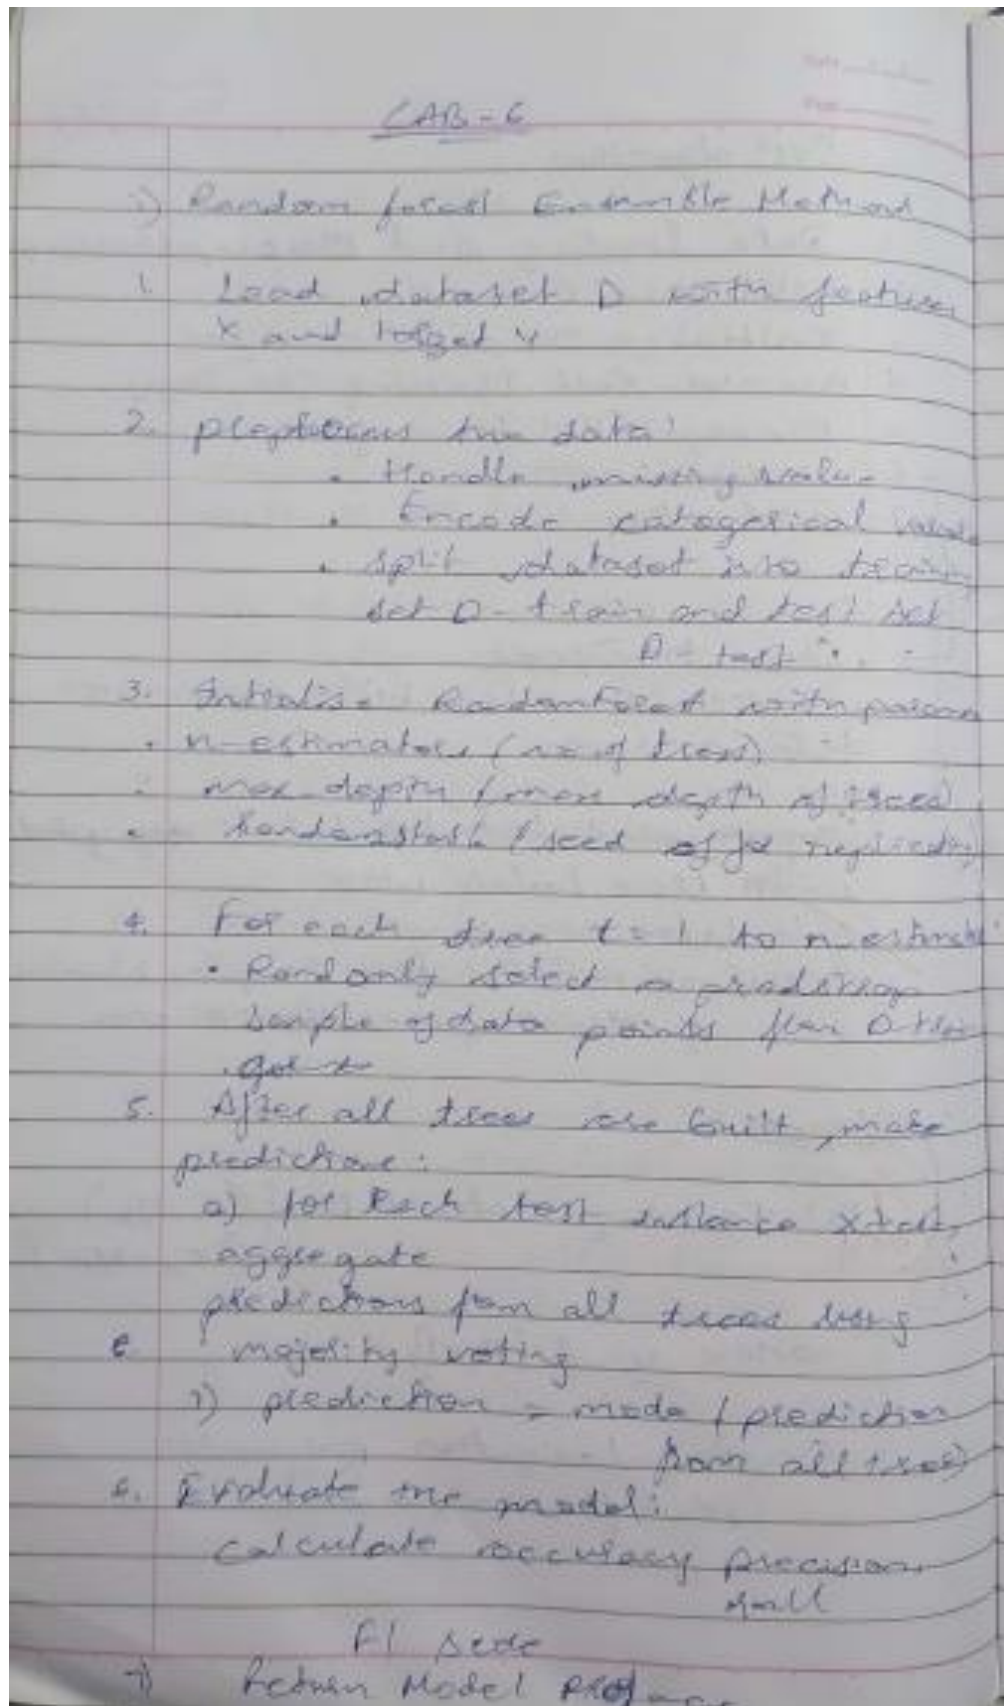
print(f"AUC score for class 0: {roc_auc[0]}")

```

## Program 8

Implement Random forest ensemble method on a given dataset.

Screenshot:



Code:

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score

data = pd.read_csv('iris.csv')
X = data.drop('species', axis=1)

y = data['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

rf_classifier = RandomForestClassifier(random_state=42)

rf_classifier.fit(X_train, y_train)

y_pred = rf_classifier.predict(X_test)

default_accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy with default n_estimators (10): {default_accuracy}")


best_n_estimators = 10

best_accuracy = default_accuracy

for n_estimators in range(1, 101):

    rf_classifier = RandomForestClassifier(n_estimators=n_estimators,
random_state=42)

    rf_classifier.fit(X_train, y_train)

    y_pred = rf_classifier.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)

    if accuracy > best_accuracy:

        best_accuracy = accuracy

        best_n_estimators = n_estimators
```



```

print(f'Best accuracy: {best_accuracy} achieved with n_estimators:
{best_n_estimators}')

from sklearn.metrics import precision_score, recall_score, f1_score

best_rf_classifier = RandomForestClassifier(n_estimators=best_n_estimators,
random_state=42)

best_rf_classifier.fit(X_train, y_train)

best_y_pred = best_rf_classifier.predict(X_test)

precision = precision_score(y_test, best_y_pred, average='weighted')

recall = recall_score(y_test, best_y_pred, average='weighted')

f1 = f1_score(y_test, best_y_pred, average='weighted')

print(f'Precision: {precision}')

print(f'Recall: {recall}')

print(f'F1-score: {f1}')

from sklearn.metrics import confusion_matrix

import seaborn as sns

import matplotlib.pyplot as plt

cm = confusion_matrix(y_test, best_y_pred)

plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',

            xticklabels=data['species'].unique(),

            yticklabels=data['species'].unique())

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

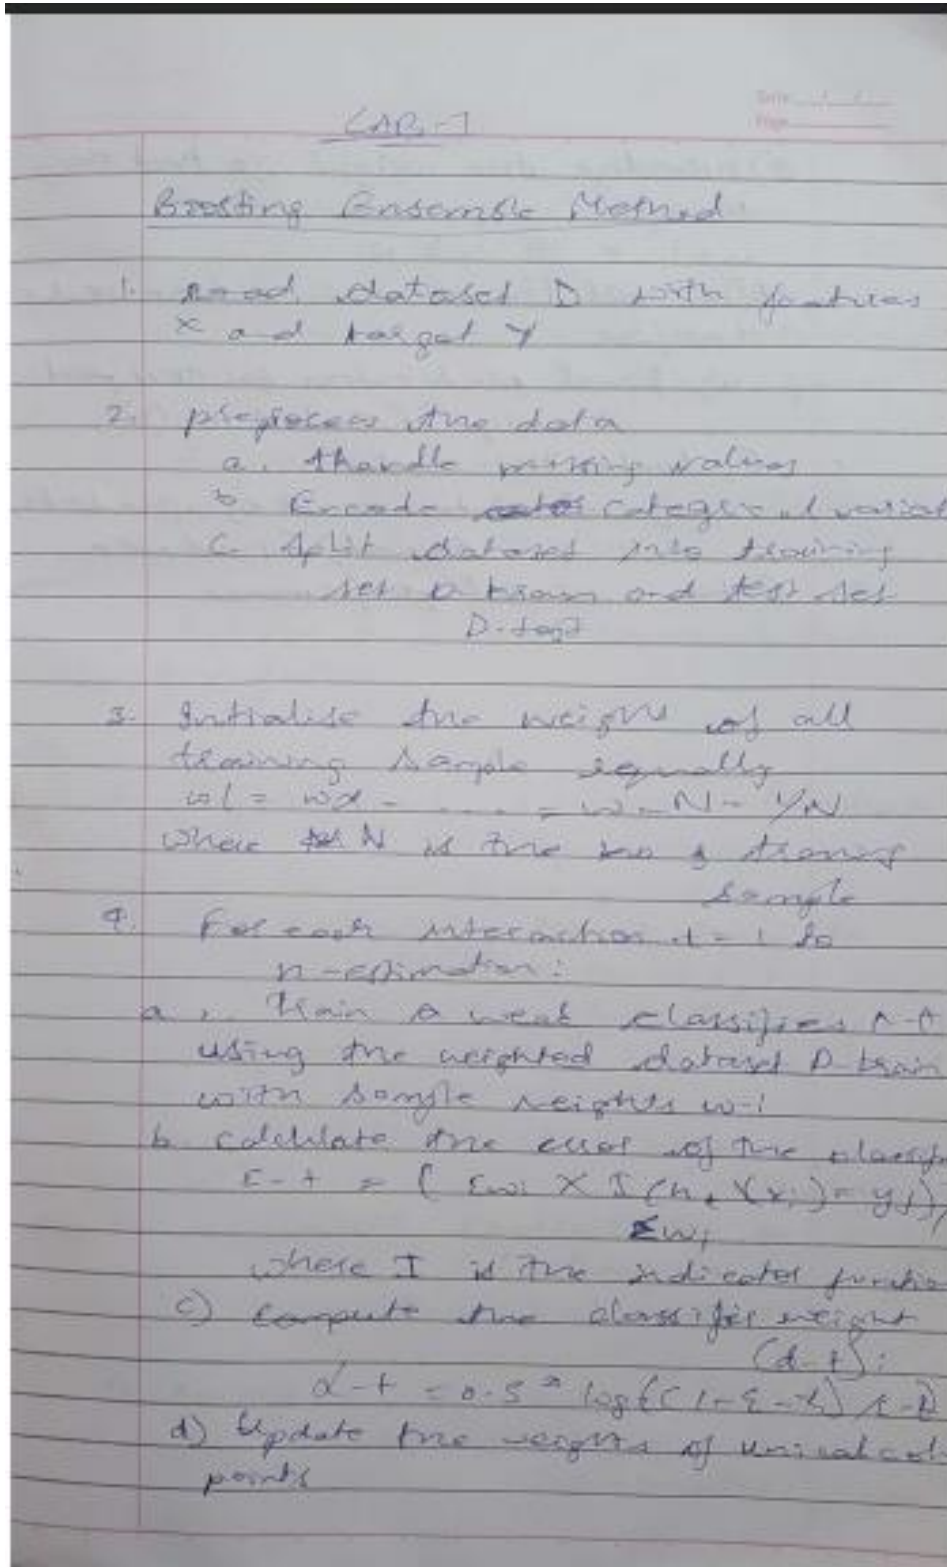
plt.show()

```

## Program 9

Implement Boosting ensemble method on a given dataset.

Screenshot:



Code:

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.ensemble import AdaBoostClassifier

from sklearn.metrics import accuracy_score

from sklearn.preprocessing import LabelEncoder

df = pd.read_csv('income.csv')

label_encoders = {}

for column in df.columns:

    if df[column].dtype == 'object':

        le = LabelEncoder()

        df[column] = le.fit_transform(df[column])
        label_encoders[column] = le

X = df.drop('income_level', axis=1)

y = df['income_level']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

abc = AdaBoostClassifier(n_estimators=10, random_state=0)

abc.fit(X_train, y_train)

y_pred = abc.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy with n_estimators=10: {accuracy}")
```

```

best_accuracy = 0

best_n_estimators = 0

for n_estimators in range(1, 101):

    abc = AdaBoostClassifier(n_estimators=n_estimators, random_state=0)

    abc.fit(X_train, y_train)

    y_pred = abc.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)

    if accuracy > best_accuracy:

        best_accuracy = accuracy

        best_n_estimators = n_estimators

print(f"\nBest accuracy: {best_accuracy} achieved with
n_estimators={best_n_estimators}")

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

import seaborn as sns

import matplotlib.pyplot as plt


abc = AdaBoostClassifier(n_estimators=73, random_state=0)

abc.fit(X_train, y_train)

y_pred = abc.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy with n_estimators=73: {accuracy}")

print(classification_report(y_test, y_pred))

cm = confusion_matrix(y_test, y_pred)

print("Confusion Matrix:")

print(cm)

plt.figure(figsize=(8, 6))

```

```
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
xticklabels=["<=50K", ">50K"], yticklabels=["<=50K", ">50K"])

plt.xlabel("Predicted")

plt.ylabel("Actual")

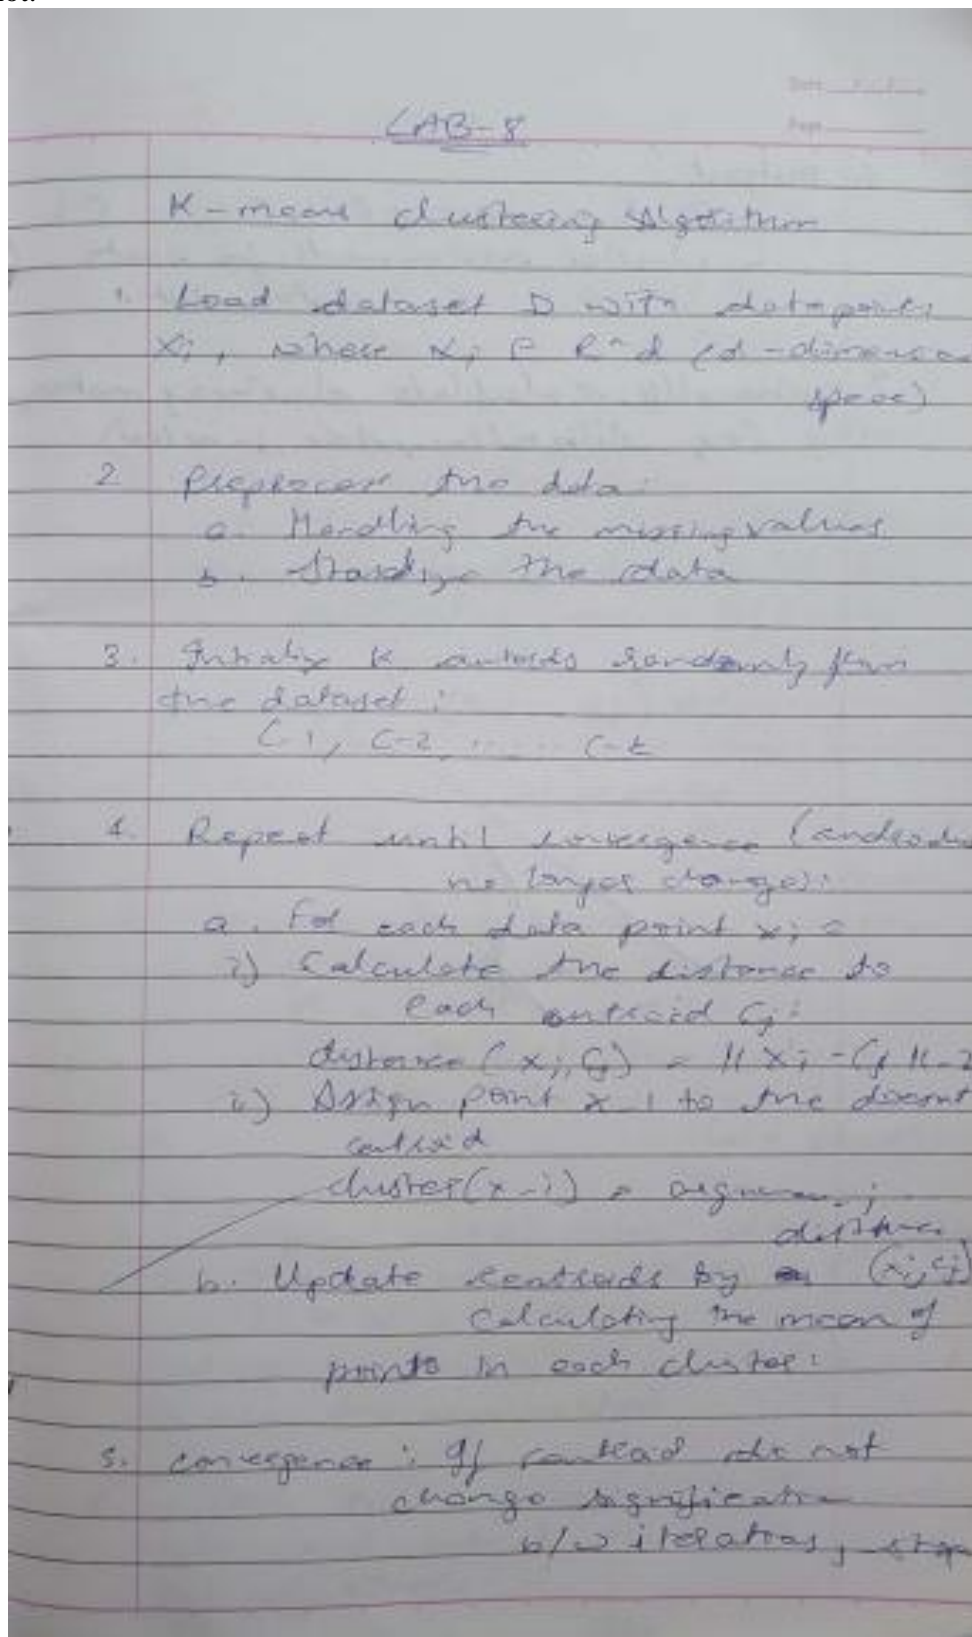
plt.title("Confusion Matrix Heatmap")

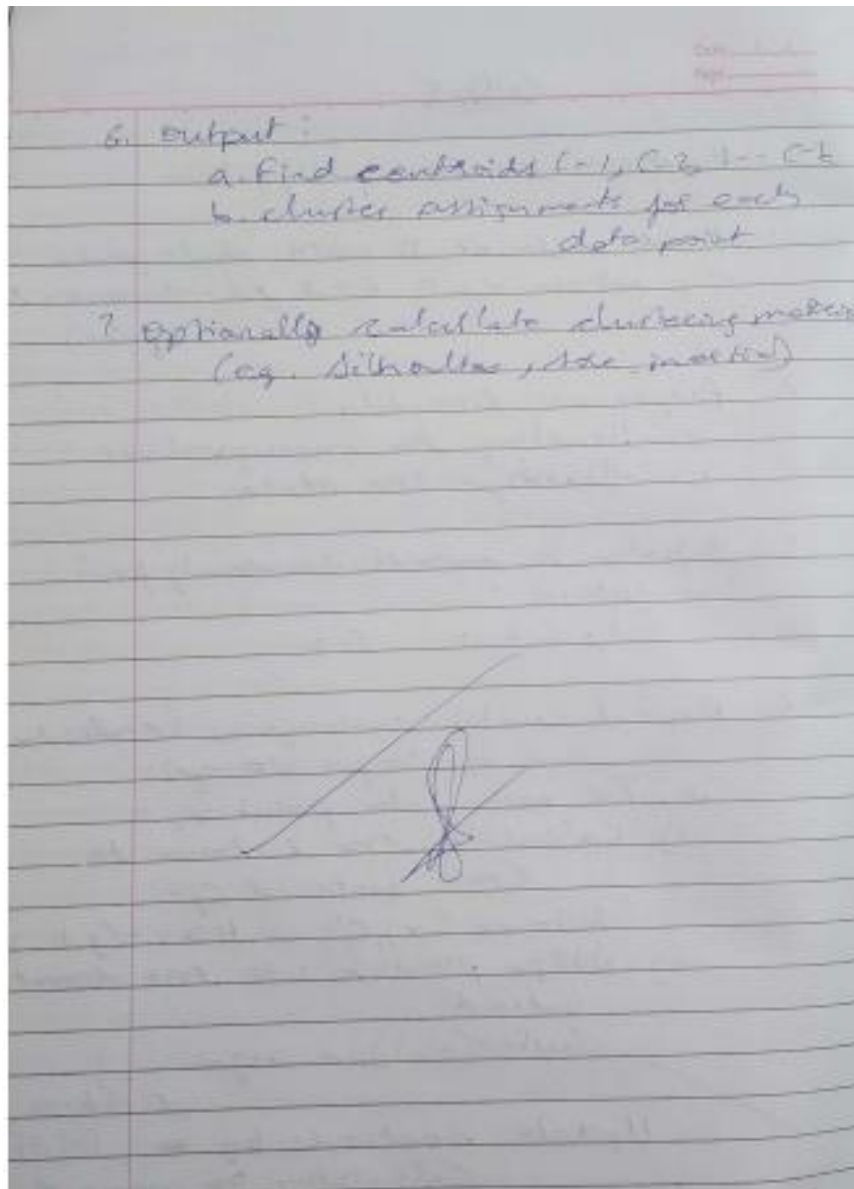
plt.show()
```

## Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Screenshot:





Code:

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.cluster import KMeans
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
df = pd.read_csv('iris.csv')
```

```
df = df[['petal_width', 'petal_length']]
```

```
scaler = MinMaxScaler()
```

```
df[['petal_width', 'petal_length']] = scaler.fit_transform(df[['petal_width',
```

```

'petal_length']])

sse = []

k_rng = range(1, 10)

for k in k_rng:

    km = KMeans(n_clusters=k)

    km.fit(df)

    sse.append(km.inertia_)


plt.xlabel('K')

plt.ylabel('Sum of squared error')

plt.plot(k_rng, sse)

plt.show()

kmeans = KMeans(n_clusters=3, random_state=0)

kmeans.fit(df)

df['cluster'] = kmeans.labels_

plt.scatter(df['petal_width'], df['petal_length'], c=df['cluster'],

cmap='viridis')

plt.xlabel('Petal Width')

plt.ylabel('Petal Length')

plt.title('K-Means Clustering of Iris Flowers')

plt.show()

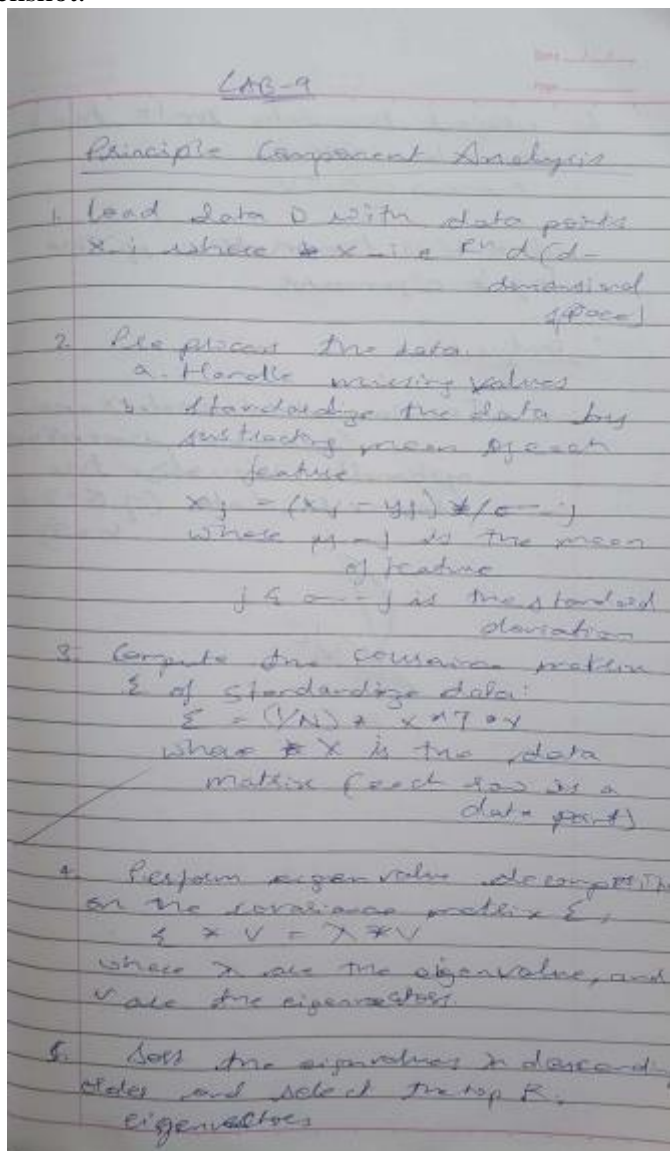
```



## Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA) method.

Screenshot:



6. project the data onto the  $k$  eigenvalues (PCA):  
 $X_{\text{new}} = X \cdot V$

where  $V$  is the matrix of the top  $k$  eigenvalues

• output

Transformed dataset  $X_{\text{new}}$   
(reduced dimension)  
optionally, visualize the  
reduced data (if  $K=2$  or  
 $K=3$ )



Code:

```
import pandas as pd

from sklearn.preprocessing import LabelEncoder, OneHotEncoder

df = pd.read_csv("heart.csv")

text_cols = df.select_dtypes(include=['object']).columns

le = LabelEncoder()

ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
for col in text_cols:

    df[col + '_le'] = le.fit_transform(df[col])

    ohe_results = ohe.fit_transform(df[[col]])

    df_ohe = pd.DataFrame(ohe_results, columns=[f"{col}_{i}" for i in
range(ohe_results.shape[1])])

    df = pd.concat([df, df_ohe], axis = 1)

df = df.drop(text_cols, axis=1)

print(df.head())

from sklearn.preprocessing import LabelEncoder, OneHotEncoder, MinMaxScaler

df = pd.read_csv("heart.csv")

text_cols = df.select_dtypes(include=['object']).columns

le = LabelEncoder()

ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False)

for col in text_cols:

    df[col + '_le'] = le.fit_transform(df[col])

    ohe_results = ohe.fit_transform(df[[col]])

    df_ohe = pd.DataFrame(ohe_results, columns=[f"{col}_{i}" for i in
```

```

range(ohe_results.shape[1]))

df = pd.concat([df, df_ohe], axis = 1)

df = df.drop(text_cols, axis=1)

scaler = MinMaxScaler()

scaled_values = scaler.fit_transform(df)

df_scaled = pd.DataFrame(scaled_values, columns=df.columns)

print(df_scaled.head())

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.linear_model import LogisticRegression

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score

X = df_scaled.drop('target', axis=1)

y = df_scaled['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

svm_model = SVC()

lr_model = LogisticRegression()

rf_model = RandomForestClassifier()

models = {'SVM': svm_model,
'Logistic Regression': lr_model,
'Random Forest': rf_model
}

results = { }

for name, model in models.items():

    model.fit(X_train, y_train)

```

```

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

results[name] = accuracy

best_model = max(results, key=results.get)

best_accuracy = results[best_model]

print("Model Accuracies:")

for name, accuracy in results.items():

    print(f"{name}: {accuracy}")

print(f"\nBest Model: {best_model} with accuracy: {best_accuracy}")

from sklearn.decomposition import PCA

pca = PCA(n_components=10)

X_train_pca = pca.fit_transform(X_train)

X_test_pca = pca.transform(X_test)

results_pca = {}

for name, model in models.items():

    model.fit(X_train_pca, y_train)

    y_pred_pca = model.predict(X_test_pca)

    accuracy_pca = accuracy_score(y_test, y_pred_pca)

    results_pca[name] = accuracy_pca

best_model_pca = max(results_pca, key=results_pca.get)

best_accuracy_pca = results_pca[best_model_pca]

print("\nModel Accuracies after PCA:")

for name, accuracy in results_pca.items():

    print(f"{name}: {accuracy}")

print(f"\nBest Model after PCA: {best_model_pca} with accuracy:

{best_accuracy_pca}")

```